

Homework 3

Kiran Charangat {style='background-color: yellow;'}

Table of contents

.....	3
Question 1	3
Question 2	8
Question 3	11
Question 4	19
Appendix	24

! Important

Please read the instructions carefully before submitting your assignment.

1. This assignment requires you to only upload a PDF file on Canvas
2. Don't collapse any code cells before submitting.
3. Remember to make sure all your code output is rendered properly before uploading your submission.

Please add your name to the author information in the frontmatter before submitting your assignment

For this assignment, we will be using the [Wine Quality](#) dataset from the UCI Machine Learning Repository. The dataset consists of red and white *vinho verde* wine samples, from the north of Portugal. The goal is to model wine quality based on physicochemical tests

We will be using the following libraries:

```
library(readr)
library(tidyr)
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(purrr)
library(car)
```

Loading required package: carData

Attaching package: 'car'

The following object is masked from 'package:purrr':

some

The following object is masked from 'package:dplyr':

recode

```
library(glmnet)
```

Loading required package: Matrix

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Loaded glmnet 4.1-8


```
make_formula <- function(x){
  as.formula(
    paste("quality ~ ", paste(x, collapse = " + "))
  )
}
```

```
# For example the following code will
# result in a formula object
# "quality ~ a + b + c"
make_formula(c("a", "b", "c"))
```

```
quality ~ a + b + c
<environment: 0x000002c7c4f83800>
```

```
make_model_matrix <- function(formula){
  X <- model.matrix(formula, df)[, -1]
  cnames <- colnames(X)
  for(i in 1:ncol(X)){
    if(!cnames[i] == "typewhite"){
      X[, i] <- scale(X[, i])
    } else {
      colnames(X)[i] <- "type"
    }
  }
  return(X)
}
```

Question 1

 50 points

Regression with categorical covariate and *t*-Test

1.1 (5 points)

Read the wine quality datasets from the specified URLs and store them in data frames **df1** and **df2**.

```
url1 <- "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv"
url2 <- "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv"

df1 <- read_delim(url1, delim = ";", show_col_types = FALSE)
df2 <- read_delim(url2, delim = ";", show_col_types = FALSE)

colnames(df1)
```

```
[1] "fixed acidity"      "volatile acidity"    "citric acid"
[4] "residual sugar"     "chlorides"           "free sulfur dioxide"
[7] "total sulfur dioxide" "density"             "pH"
[10] "sulphates"          "alcohol"             "quality"
```

```
head(df1)
```

```
# A tibble: 6 x 12
  `fixed acidity` `volatile acidity` `citric acid` `residual sugar` chlorides
      <dbl>          <dbl>          <dbl>          <dbl>      <dbl>
1           7          0.27          0.36          20.7      0.045
2          6.3          0.3          0.34           1.6      0.049
3           8.1          0.28          0.4           6.9      0.05
4           7.2          0.23          0.32           8.5      0.058
5           7.2          0.23          0.32           8.5      0.058
6           8.1          0.28          0.4           6.9      0.05
# i 7 more variables: `free sulfur dioxide` <dbl>,
#   `total sulfur dioxide` <dbl>, density <dbl>, pH <dbl>, sulphates <dbl>,
#   alcohol <dbl>, quality <dbl>
```

1.2 (5 points)

Perform the following tasks to prepare the data frame `df` for analysis:

1. Combine the two data frames into a single data frame `df`, adding a new column called `type` to indicate whether each row corresponds to white or red wine.
2. Rename the columns of `df` to replace spaces with underscores
3. Remove the columns `fixed_acidity` and `free_sulfur_dioxide`

4. Convert the `type` column to a factor
5. Remove rows (if any) with missing values.

```
df <- bind_rows(df1 %>% mutate(type = "white"),
                df2 %>% mutate(type = "red")) %>%
  mutate(type = as.factor(type)) %>%
  drop_na() %>%
  rename_with(~ gsub(" ", "_", .x)) %>%
  select(-fixed_acidity, -free_sulfur_dioxide)
```

```
colnames(df)
```

```
[1] "volatile_acidity"      "citric_acid"          "residual_sugar"
[4] "chlorides"            "total_sulfur_dioxide" "density"
[7] "pH"                   "sulphates"            "alcohol"
[10] "quality"              "type"
```

```
dim(df)
```

```
[1] 6497  11
```

Your output to R `dim(df)` should be

```
[1] 6497  11
```

1.3 (20 points)

Recall from STAT 200, the method to compute the t statistic for the the difference in means (with the equal variance assumption)

1. Using `df` compute the mean of `quality` for red and white wine separately, and then store the difference in means as a variable called `diff_mean`.
2. Compute the pooled sample variance and store the value as a variable called `sp_squared`.
3. Using `sp_squared` and `diff_mean`, compute the t Statistic, and store its value in a variable called `t1`.

```
diff_mean <-df %>%
  group_by(type) %>%
  summarise(mean_quality = mean(quality)) %>%
  summarise(diff_mean = diff(mean_quality)) %>%
  pull(diff_mean)

diff_mean
```

```
[1] 0.2418868
```

```
sp_squared <- df %>%
  group_by(type) %>%
  summarise(n = n(),
            var_quality = var(quality)) %>%
  summarise(sp_squared = ((n[1] - 1) * var_quality[1] + (n[2] - 1) *
  pull(sp_squared)

sp_squared
```

```
[1] 0.7518329
```

```
t1 <- df %>%
  summarise(t1 = diff_mean / sqrt(sp_squared * (1/sum(df$type == "white") + 1
  pull(t1)

t1
```

```
[1] 9.68565
```

1.4 (10 points)

Equivalently, R has a function called `t.test()` which enables you to perform a two-sample *t*-Test without having to compute the pooled variance and difference in means.

Perform a two-sample *t*-test to compare the quality of white and red wines using the `t.test()` function with the setting `var.equal=TRUE`. Store the *t*-statistic in `t2`.

```
t_test <- t.test(quality ~ type, data = df, var.equal = TRUE)
t_test
```

Two Sample t-test

```
data: quality by type
t = -9.6856, df = 6495, p-value < 2.2e-16
alternative hypothesis: true difference in means between group red and group white is not equal to 0
95 percent confidence interval:
 -0.2908436 -0.1929301
sample estimates:
mean in group red mean in group white
      5.636023      5.877909
```

```
t2 <- t_test$statistic
t2
```

```
t
-9.68565
```

1.5 (5 points)

Fit a linear regression model to predict `quality` from `type` using the `lm()` function, and extract the t -statistic for the `type` coefficient from the model summary. Store this t -statistic in `t3`.

```
fit <- lm(quality ~ type, data = df)
fit
```

Call:

```
lm(formula = quality ~ type, data = df)
```

Coefficients:

```
(Intercept)      typewhite
      5.6360      0.2419
```

```
t3 <- summary(fit)$coefficients[2, "t value"]  
t3
```

[1] 9.68565

1.6 (5 points)

Print a vector containing the values of `t1`, `t2`, and `t3`. What can you conclude from this? Why?

```
t_statistics <- c(t1, t2, t3)  
t_statistics
```

```
           t  
9.68565 -9.68565  9.68565
```

```
# These t-statistic values indicate a significant difference in wine quality  
# between red and white wine. The negative sign in the second value (t2) suggests  
# a reversal in direction of the effect. These results, however, show that the  
# wine type significantly impacts quality, but it is important to understand the  
# direction of the effect.
```

Question 2

💡 25 points

Collinearity

2.1 (5 points)

Fit a linear regression model with all predictors against the response variable `quality`. Use the `broom::tidy()` function to print a summary of the fitted model. What can we conclude from the model summary?


```
broom::tidy(fit)
```

```
# A tibble: 2 x 5
```

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	5.64	0.0217	260.	0
2	typewhite	0.242	0.0250	9.69	4.89e-22

```
# The coefficient for 'typewhite' indicates the estimated change in quality score when com
```

2.2 (10 points)

Fit two **simple** linear regression models using `lm()`: one with only `citric_acid` as the predictor, and another with only `total_sulfur_dioxide` as the predictor. In both models, use `quality` as the response variable. How does your model summary compare to the summary from the previous question?

```
model_citric <- lm(quality ~ citric_acid, data = df)
```

```
model_sulfur <- lm(quality ~ total_sulfur_dioxide, data = df)
```


2.3 (5 points)

Visualize the correlation matrix of all numeric columns in `df` using `corrplot()`

```
library(corrplot)
```

corrplot 0.92 loaded

```
df %>%  
  select_if(is.numeric) %>%  
  cor() %>%  
  corrplot(method = "circle")
```



2.4 (5 points)

Compute the variance inflation factor (VIF) for each predictor in the full model using `vif()` function. What can we conclude from this?

```
vif(lm(quality ~ ., data = df))
```

volatile_acidity	citric_acid	residual_sugar
2.103853	1.549248	4.680035
chlorides	total_sulfur_dioxide	density
1.625065	2.628534	9.339357
pH	sulphates	alcohol
1.352005	1.522809	3.419849
type		
6.694679		

The predictor density has a VIF of over 9, and type has a VIF of close to 7, which suggests

Question 3

💡 40 points

Variable selection

3.1 (5 points)

Run a backward stepwise regression using a `full_model` object as the starting model. Store the final formula in an object called `backward_formula` using the built-in `formula()` function in R

```
full_model <- lm(quality ~ ., data = df)
backward_model <- step(full_model, direction = "backward")
```

Start: AIC=-3953.43

```
quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
  total_sulfur_dioxide + density + pH + sulphates + alcohol +
  type
```

	Df	Sum of Sq	RSS	AIC
- citric_acid	1	0.066	3523.6	-3955.3
- total_sulfur_dioxide	1	0.854	3524.4	-3953.9
<none>			3523.5	-3953.4
- pH	1	4.413	3527.9	-3947.3
- chlorides	1	4.559	3528.1	-3947.0
- density	1	19.054	3542.6	-3920.4
- type	1	26.794	3550.3	-3906.2
- sulphates	1	41.399	3564.9	-3879.5
- residual_sugar	1	63.881	3587.4	-3838.7
- alcohol	1	206.860	3730.4	-3584.8
- volatile_acidity	1	216.549	3740.0	-3567.9

Step: AIC=-3955.3

```
quality ~ volatile_acidity + residual_sugar + chlorides + total_sulfur_dioxide +
  density + pH + sulphates + alcohol + type
```

	Df	Sum of Sq	RSS	AIC
- total_sulfur_dioxide	1	0.818	3524.4	-3955.8

<none>			3523.6	-3955.3
- chlorides	1	4.495	3528.1	-3949.0
- pH	1	4.536	3528.1	-3948.9
- density	1	20.794	3544.4	-3919.1
- type	1	26.943	3550.5	-3907.8
- sulphates	1	41.491	3565.1	-3881.2
- residual_sugar	1	67.371	3590.9	-3834.3
- alcohol	1	235.151	3758.7	-3537.6
- volatile_acidity	1	252.565	3776.1	-3507.5

Step: AIC=-3955.8

quality ~ volatile_acidity + residual_sugar + chlorides + density +
pH + sulphates + alcohol + type

	Df	Sum of Sq	RSS	AIC
<none>			3524.4	-3955.8
- pH	1	4.295	3528.7	-3949.9
- chlorides	1	4.523	3528.9	-3949.5
- density	1	21.540	3545.9	-3918.2
- sulphates	1	40.711	3565.1	-3883.2
- type	1	43.664	3568.0	-3877.8
- residual_sugar	1	66.572	3591.0	-3836.2
- alcohol	1	244.545	3768.9	-3521.9
- volatile_acidity	1	256.695	3781.1	-3501.0

```
backward_formula <- formula(backward_model)
backward_formula
```

quality ~ volatile_acidity + residual_sugar + chlorides + density +
pH + sulphates + alcohol + type

3.2 (5 points)

Run a forward stepwise regression using a `null_model` object as the starting model. Store the final formula in an object called `forward_formula` using the built-in `formula()` function in R

```
null_model <- lm(quality ~ 1, data = df)
forward_model <- step(null_model, direction = "forward", scope = formula(full_model))
```

Start: AIC=-1760.04
quality ~ 1

	Df	Sum of Sq	RSS	AIC
+ alcohol	1	977.95	3975.7	-3186.9
+ density	1	463.41	4490.3	-2396.2
+ volatile_acidity	1	349.71	4604.0	-2233.7
+ chlorides	1	199.47	4754.2	-2025.1
+ type	1	70.53	4883.2	-1851.2
+ citric_acid	1	36.24	4917.4	-1805.7
+ total_sulfur_dioxide	1	8.48	4945.2	-1769.2
+ sulphates	1	7.34	4946.3	-1767.7
+ residual_sugar	1	6.77	4946.9	-1766.9
+ pH	1	1.88	4951.8	-1760.5
<none>			4953.7	-1760.0

Step: AIC=-3186.88
quality ~ alcohol

	Df	Sum of Sq	RSS	AIC
+ volatile_acidity	1	307.508	3668.2	-3707.9
+ residual_sugar	1	85.662	3890.1	-3326.4
+ type	1	54.335	3921.4	-3274.3
+ citric_acid	1	40.303	3935.4	-3251.1
+ chlorides	1	39.696	3936.0	-3250.1
+ total_sulfur_dioxide	1	31.346	3944.4	-3236.3
+ sulphates	1	7.859	3967.9	-3197.7
+ pH	1	5.938	3969.8	-3194.6
<none>			3975.7	-3186.9
+ density	1	0.005	3975.7	-3184.9

Step: AIC=-3707.89
quality ~ alcohol + volatile_acidity

	Df	Sum of Sq	RSS	AIC
+ sulphates	1	48.259	3620.0	-3791.9
+ density	1	38.704	3629.5	-3774.8
+ residual_sugar	1	29.751	3638.5	-3758.8
+ type	1	28.895	3639.3	-3757.3
+ total_sulfur_dioxide	1	5.619	3662.6	-3715.9
+ pH	1	5.533	3662.7	-3715.7
<none>			3668.2	-3707.9
+ chlorides	1	0.162	3668.1	-3706.2

+ citric_acid	1	0.099	3668.1	-3706.1
---------------	---	-------	--------	---------

Step: AIC=-3791.94

quality ~ alcohol + volatile_acidity + sulphates

	Df	Sum of Sq	RSS	AIC
+ residual_sugar	1	43.989	3576.0	-3869.4
+ density	1	18.661	3601.3	-3823.5
+ type	1	6.012	3614.0	-3800.7
+ chlorides	1	4.988	3615.0	-3798.9
+ citric_acid	1	2.031	3617.9	-3793.6
+ pH	1	1.903	3618.1	-3793.4
<none>			3620.0	-3791.9
+ total_sulfur_dioxide	1	0.817	3619.2	-3791.4

Step: AIC=-3869.37

quality ~ alcohol + volatile_acidity + sulphates + residual_sugar

	Df	Sum of Sq	RSS	AIC
+ type	1	20.7581	3555.2	-3905.2
+ total_sulfur_dioxide	1	13.3542	3562.6	-3891.7
+ pH	1	6.6430	3569.3	-3879.5
+ citric_acid	1	4.3384	3571.6	-3875.3
+ chlorides	1	1.8907	3574.1	-3870.8
<none>			3576.0	-3869.4
+ density	1	0.0071	3576.0	-3867.4

Step: AIC=-3905.19

quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
type

	Df	Sum of Sq	RSS	AIC
+ density	1	20.4623	3534.8	-3940.7
+ chlorides	1	6.6602	3548.6	-3915.4
+ citric_acid	1	5.2242	3550.0	-3912.7
+ pH	1	3.9477	3551.3	-3910.4
+ total_sulfur_dioxide	1	1.2539	3554.0	-3905.5
<none>			3555.2	-3905.2

Step: AIC=-3940.7

quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
type + density

	Df	Sum of Sq	RSS	AIC
+ chlorides	1	6.0826	3528.7	-3949.9
+ pH	1	5.8541	3528.9	-3949.5
<none>			3534.8	-3940.7
+ citric_acid	1	0.8471	3533.9	-3940.3
+ total_sulfur_dioxide	1	0.5646	3534.2	-3939.7

Step: AIC=-3949.89

quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
type + density + chlorides

	Df	Sum of Sq	RSS	AIC
+ pH	1	4.2945	3524.4	-3955.8
<none>			3528.7	-3949.9
+ total_sulfur_dioxide	1	0.5765	3528.1	-3948.9
+ citric_acid	1	0.2338	3528.4	-3948.3

Step: AIC=-3955.8

quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
type + density + chlorides + pH

	Df	Sum of Sq	RSS	AIC
<none>			3524.4	-3955.8
+ total_sulfur_dioxide	1	0.81762	3523.6	-3955.3
+ citric_acid	1	0.02919	3524.4	-3953.9

```
forward_formula <- formula(forward_model)
forward_formula
```

quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
type + density + chlorides + pH

3.3 (10 points)

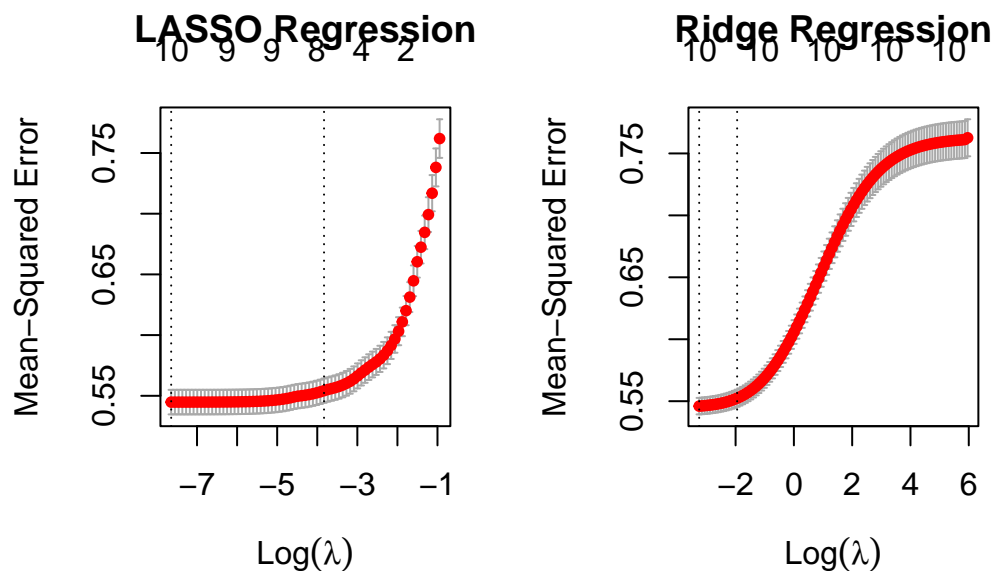
1. Create a `y` vector that contains the response variable (`quality`) from the `df` dataframe.
2. Create a design matrix `X` for the `full_model` object using the `make_model_matrix()` function provided in the Appendix.

- Then, use the `cv.glmnet()` function to perform LASSO and Ridge regression with X and y .

```
y <- df$quality
X <- make_model_matrix(formula(full_model))
lasso_model <- cv.glmnet(X, y, alpha = 1)
ridge_model <- cv.glmnet(X, y, alpha = 0)
```

Create side-by-side plots of the ridge and LASSO regression results. Interpret your main findings.

```
par(mfrow=c(1, 2))
plot(lasso_model, main="LASSO Regression")
plot(ridge_model, main="Ridge Regression")
```



From these models, it be concluded that both models benefit from regularization, as indi

3.4 (5 points)

Print the coefficient values for LASSO regression at the `lambda.1se` value? What are the variables selected by LASSO?

```
lasso_coef_at_1se <- coef(lasso_model, s = "lambda.1se")
lasso_coef_at_1se
```

11 x 1 sparse Matrix of class "dgCMatrix"

```
              s1
(Intercept)    5.822812544
volatile_acidity -0.204848960
citric_acid      .
residual_sugar   0.059537104
chlorides        -0.001916886
total_sulfur_dioxide -0.008115301
density          .
pH               0.003557628
sulphates        0.067073695
alcohol          0.377201730
type            -0.005882625
```

```
# The variables selected by LASSO are volatile_acidity with a coefficient of approximately
```

Store the variable names with non-zero coefficients in `lasso_vars`, and create a formula object called `lasso_formula` using the `make_formula()` function provided in the Appendix.

```
lasso_coef_vector <- as.numeric(lasso_coef_at_1se)
names(lasso_coef_vector) <- rownames(lasso_coef_at_1se)
lasso_vars <- names(lasso_coef_vector)[-1][lasso_coef_vector[-1] != 0]
lasso_formula <- make_formula(lasso_vars)
lasso_formula
```

```
quality ~ volatile_acidity + residual_sugar + chlorides + total_sulfur_dioxide +
  pH + sulphates + alcohol + type
<environment: 0x000002c7ccb2bd20>
```

3.5 (5 points)

Print the coefficient values for ridge regression at the `lambda.1se` value? What are the variables selected here?

```
ridge_coef_at_1se <- coef(ridge_model, s = "lambda.1se")
ridge_coef_at_1se
```

11 x 1 sparse Matrix of class "dgCMatrix"

	s1
(Intercept)	5.89912780
volatile_acidity	-0.19275639
citric_acid	0.01625926
residual_sugar	0.11607371
chlorides	-0.04421587
total_sulfur_dioxide	-0.04011160
density	-0.09019465
pH	0.02601554
sulphates	0.08689460
alcohol	0.28992182
type	-0.10711174

Store the variable names with non-zero coefficients in `ridge_vars`, and create a formula object called `ridge_formula` using the `make_formula()` function provided in the Appendix.

```
ridge_vars <- rownames(ridge_coef_at_1se)[-1]
ridge_formula <- make_formula(ridge_vars)
print(ridge_formula)
```

```
quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
  total_sulfur_dioxide + density + pH + sulphates + alcohol +
  type
<environment: 0x000002c7ce6e8e28>
```

3.6 (10 points)

What is the difference between stepwise selection, LASSO and ridge based on you analyses above?

```
# The stepwise selection would have given me a subset of predictors based on a criterion 1
# The LASSO regression results showed which variables were most influential by keeping the
# The Ridge regression results would have included all predictors but with their influence
```

Question 4

💡 70 points

Variable selection

4.1 (5 points)

Excluding `quality` from `df` we have 10 possible predictors as the covariates. How many different models can we create using any subset of these 10 covariates as possible predictors? Justify your answer.

```
# We can create  $2^{10} - 1$  which is 1023 different models. Each predictor can either be incl
```

4.2 (20 points)

Store the names of the predictor variables (all columns except `quality`) in an object called `x_vars`.

```
x_vars <- colnames(df %>% select(-quality))
```

```
x_vars <- colnames(df %>% select(-quality))
```

Use:

- the `combn()` function (built-in R function) and
- the `make_formula()` (provided in the Appendix)

to **generate all possible linear regression formulas** using the variables in `x_vars`. This is most optimally achieved using the `map()` function from the `purrr` package.

```
formulas <- map(
  1:length(x_vars),
  \(x){
    vars <- combn(x_vars, x, simplify = FALSE)
    map(vars, \(combo) make_formula(combo))
  }
) %>% unlist(recursive = FALSE)
```

If your code is right the following command should return something along the lines of:

```
sample(formulas, 4) %>% as.character()
```

```
[1] "quality ~ volatile_acidity + citric_acid + chlorides + pH + type"
[2] "quality ~ citric_acid + residual_sugar + chlorides + total_sulfur_dioxide + sulphates +
[3] "quality ~ residual_sugar + chlorides + total_sulfur_dioxide + density"
[4] "quality ~ volatile_acidity + citric_acid + density"
```

```
sample(formulas, 4) %>% as.character()
# Output:
# [1] "quality ~ volatile_acidity + residual_sugar + density + pH + alcohol"
# [2] "quality ~ citric_acid"
# [3] "quality ~ volatile_acidity + citric_acid + residual_sugar + total_sulfur_dioxide +
# [4] "quality ~ citric_acid + chlorides + total_sulfur_dioxide + pH + alcohol + type"
```

4.3 (10 points)

Use `map()` and `lm()` to fit a linear regression model to each formula in `formulas`, using `df` as the data source. Use `broom::glance()` to extract the model summary statistics, and bind them together into a single tibble of summaries using the `bind_rows()` function from `dplyr`.

```
models <- map(formulas, ~lm(.x, data = df))
summaries <- map(models, ~broom::glance(.x))
all_summaries <- bind_rows(summaries)
all_summaries
```

A tibble: 1,023 x 12

	r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.0706	0.0705	0.842	493.	2.06e-105	1	-8100.	16206.	16226.
2	0.00732	0.00716	0.870	47.9	5.00e-12	1	-8314.	16634.	16654.
3	0.00137	0.00121	0.873	8.89	2.87e-3	1	-8333.	16673.	16693.
4	0.0403	0.0401	0.856	273.	5.32e-60	1	-8204.	16415.	16435.
5	0.00171	0.00156	0.873	11.1	8.48e-4	1	-8332.	16671.	16691.
6	0.0935	0.0934	0.831	670.	9.66e-141	1	-8019.	16044.	16064.
7	0.000380	0.000227	0.873	2.47	1.16e-1	1	-8337.	16679.	16700.
8	0.00148	0.00133	0.873	9.63	1.92e-3	1	-8333.	16672.	16692.

```
9 0.197      0.197    0.782   1598.   1.50e-312    1 -7623. 15253. 15273.
10 0.0142     0.0141   0.867    93.8   4.89e- 22    1 -8291. 16588. 16609.
# i 1,013 more rows
# i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

4.4 (5 points)

Extract the `adj.r.squared` values from `summaries` and use them to identify the formula with the *highest* adjusted R-squared value.

```
best_model_index <- which.max(all_summaries$adj.r.squared)
rsq_formula <- formulas[[best_model_index]]
rsq_formula
```

```
quality ~ volatile_acidity + residual_sugar + chlorides + total_sulfur_dioxide +
  density + pH + sulphates + alcohol + type
<environment: 0x000002c7cfbf2810>
```

Store resulting formula as a variable called `rsq_formula`.

4.5 (5 points)

Extract the AIC values from `summaries` and use them to identify the formula with the *lowest* AIC value.

```
lowest_aic_index <- which.min(all_summaries$AIC)
aic_formula <- formulas[[lowest_aic_index]]
aic_formula
```

```
quality ~ volatile_acidity + residual_sugar + chlorides + density +
  pH + sulphates + alcohol + type
<environment: 0x000002c7cfbb89c8>
```

Store resulting formula as a variable called `aic_formula`.

4.6 (15 points)

Combine all formulas shortlisted into a single vector called `final_formulas`.

```
null_formula <- formula(null_model)
full_formula <- formula(full_model)

final_formulas <- c(
  null_formula,
  full_formula,
  backward_formula,
  forward_formula,
  lasso_formula,
  ridge_formula,
  rsq_formula,
  aic_formula
)
```

- Are `aic_formula` and `rsq_formula` the same? How do they differ from the formulas shortlisted in question 3?
 - The `aic_formula` and `rsq_formula` are almost similar as the `rsq_formula` includes `total_sulfur_dioxide` and `aic_formula` does not. As for how they differ from the formulas shortlisted in question 3, The `aic_formula` and `rsq_formula` represent models selected based on optimizing specific statistical criteria. The formulas from question 3 however, like the `ridge_formula` and `lasso_formula` apply regularization techniques and the forward and backward formulas use step wise procedures.
- Which of these is more reliable? Why?
 - For prediction purposes and generalization, `aic_formula` and `lasso_formula` might be good options. For explanatory purposes, `rsq_formula` might be good like when explaining the variance in response is the primary goal.
- If we had a dataset with 10,000 columns, which of these methods would you consider for your analyses? Why?
 - I would choose the LASSO and Ridge regression for this analysis due to its effectiveness in handling multi-collinearity and its ability to prevent over fitting through regularization.

4.7 (10 points)

Use `map()` and `glance()` to extract the `sigma`, `adj.r.squared`, `AIC`, `df`, and `p.value` statistics for each model obtained from `final_formulas`. Bind them together into a single data frame `summary_table`. Summarize your main findings.

```
summary_table <- map(
  final_formulas,
  \(x) {
    model <- lm(x, data = df)
    broom::glance(model) %>%
      select(sigma, adj.r.squared, AIC, df = df.residual, p.value = p.value)
  }
) %>% bind_rows()

summary_table %>% knitr::kable()
```

sigma	adj.r.squared	AIC	df	p.value
0.8732553	0.0000000	16679.64	6496	NA
0.7370527	0.2876152	14486.26	6486	0
0.7370314	0.2876563	14483.89	6488	0
0.7370314	0.2876563	14483.89	6488	0
0.7391172	0.2836187	14520.61	6488	0
0.7370527	0.2876152	14486.26	6486	0
0.7370027	0.2877118	14484.38	6487	0
0.7370314	0.2876563	14483.89	6488	0

Appendix

Convenience function for creating a formula object

The following function which takes as input a vector of column names `x` and outputs a **formula** object with `quality` as the response variable and the columns of `x` as the covariates.

```
make_formula <- function(x){  
  as.formula(  
    paste("quality ~ ", paste(x, collapse = " + "))  
  )  
}  
  
# For example the following code will  
# result in a formula object  
# "quality ~ a + b + c"  
make_formula(c("a", "b", "c"))
```

```
quality ~ a + b + c  
<environment: 0x000002c7889442c0>
```

Convenience function for glmnet

The `make_model_matrix` function below takes a **formula** as input and outputs a **rescaled** model matrix `X` in a format amenable for `glmnet()`

```
make_model_matrix <- function(formula){  
  X <- model.matrix(formula, df)[, -1]  
  cnames <- colnames(X)  
  for(i in 1:ncol(X)){  
    if(!cnames[i] == "typewhite"){  
      X[, i] <- scale(X[, i])  
    } else {  
      colnames(X)[i] <- "type"  
    }  
  }  
  return(X)  
}
```


Session Information

Print your R session information using the following command

```
sessionInfo()
```

R version 4.3.2 (2023-10-31 ucrt)

Platform: x86_64-w64-mingw32/x64 (64-bit)

Running under: Windows 11 x64 (build 22621)

Matrix products: default

locale:

```
[1] LC_COLLATE=English_United States.utf8
[2] LC_CTYPE=English_United States.utf8
[3] LC_MONETARY=English_United States.utf8
[4] LC_NUMERIC=C
[5] LC_TIME=English_United States.utf8
```

time zone: America/New_York

tzcode source: internal

attached base packages:

```
[1] stats      graphics  grDevices  utils      datasets  methods    base
```

other attached packages:

```
[1] corrplot_0.92  glmnet_4.1-8   Matrix_1.6-1.1 car_3.1-2      carData_3.0-5
[6] purrr_1.0.2    dplyr_1.1.4    tidyr_1.3.0    readr_2.1.5
```

loaded via a namespace (and not attached):

```
[1] utf8_1.2.4      generics_0.1.3  shape_1.4.6     lattice_0.21-9
[5] hms_1.1.3       digest_0.6.33   magrittr_2.0.3  evaluate_0.23
[9] grid_4.3.2      iterators_1.0.14 fastmap_1.1.1    foreach_1.5.2
[13] jsonlite_1.8.7  backports_1.4.1 survival_3.5-7   fansi_1.0.5
[17] codetools_0.2-19 abind_1.4-5     cli_3.6.1       rlang_1.1.2
[21] crayon_1.5.2    bit64_4.0.5     splines_4.3.2   withr_2.5.2
[25] yaml_2.3.7      tools_4.3.2     parallel_4.3.2  tzdb_0.4.0
[29] broom_1.0.5     curl_5.1.0      vctrs_0.6.4     R6_2.5.1
[33] lifecycle_1.0.4 bit_4.0.5       vroom_1.6.4     pkgconfig_2.0.3
[37] pillar_1.9.0    glue_1.6.2      Rcpp_1.0.11     xfun_0.41
[41] tibble_3.2.1    tidyselect_1.2.0 rstudioapi_0.15.0 knitr_1.45
```

[45] `htmltools_0.5.7` `rmarkdown_2.25` `compiler_4.3.2`